

logistic-regression

May 8, 2018

1 Loading the Dataset

```
In [4]: import numpy as np
        from sklearn.datasets import load_digits
        digits = load_digits()
        print(digits.keys())
        data = digits["data"]
        images = digits["images"]
        target = digits["target"]
        target_names = digits["target_names"]
        print(data.shape)
```

```
dict_keys(['data', 'target', 'target_names', 'images', 'DESCR'])
(1797, 64)
```

Extract instances showing “3” or “8”, append a column of “1s” and create a vector of ground-truth labels where 1 corresponds to 3 and -1 to 8.

```
In [5]: X = data[np.logical_or(target == 3, target == 8)]
        y = target[np.logical_or(target == 3, target == 8)]

        X = np.concatenate((X, np.ones([len(y),1])), axis = 1)

        y[y==3] = 1
        y[y==8] = -1
```

2 1.1 Classification with sklearn

```
In [6]: from sklearn.linear_model import LogisticRegression
        from sklearn.model_selection import cross_val_score, train_test_split

In [7]: num_splits = 10
        lambdas = [0.001,0.01,0.1,1,10,100,1000]
        scores = np.zeros([len(lambdas), 2]) # save mean score and std for each lambda

        for i, C in enumerate(lambdas):
```

```

logistic = LogisticRegression(C = C)
curr_scores = cross_val_score(logistic, X, y, cv = num_splits)
scores[i,0] = np.average(curr_scores)
scores[i,1] = np.std(curr_scores)

```

```

In [8]: import pandas as pd
        pd.options.display.float_format = '{:,.4f}'.format

```

```

In [9]: display(pd.DataFrame(
        data = scores,
        index = lambdas,
        columns = ['mean', 'std'])
        .rename_axis('C', axis = 'columns'))

```

C	mean	std
0.0010	0.9688	0.0531
0.0100	0.9691	0.0464
0.1000	0.9860	0.0288
1.0000	0.9860	0.0288
10.0000	0.9803	0.0283
100.0000	0.9775	0.0304
1,000.0000	0.9747	0.0341

The accuracy of the prediction varies only very slowly when modifying the regularization parameter. The best accuracy is obtained for $\lambda = 0.1, 1, 10$, thus a very large range. In the following we will proceed using $\lambda = 1$.

3 1.2 Optimization Methods

```

In [10]: def sigmoid(z):
        return 1 / (1 + np.exp(-z))

```

```

In [11]: def predict(beta, X):
        # if X*beta > 0 --> y=1, if X*beta < 0 --> y=-1
        y = (X.dot(beta) >= 0) * 2 - 1
        return y

```

```

In [12]: def zero_one_loss(y_prediction, y_truth):
        return np.sum(np.not_equal(y_prediction, y_truth))

```

```

In [13]: def gradient(beta, X, y, lambda_ = 1):
        # distinguish the cases of one single / several training instances
        if np.isscalar(y):
            grad = beta / lambda_ - sigmoid(-X.dot(beta) * y) * y * X
        else:
            grad = beta / lambda_ - np.average((sigmoid(-X.dot(beta) * y) * y)[:,None] * X,
            return grad

```

```
In [14]: def GD(X, y, beta, tau_0, gamma, m):
        for _ in range(m):
            tau = tau_0 / (1 + gamma * _)
            beta = beta - tau * gradient(beta, X, y)
        return beta
```

To sample single random instances, we do sampling without replacement since we want to avoid computing the gradient from the same instance repetitively. Instead, we want to use as much different samples as possible to compute the gradient. Only in the case of mini-batches, we do sampling with replacement since the training set is not large enough to sample 150 times a mini-batch without replacing “old” samples.

```
In [15]: from sklearn.utils import shuffle
        def SGD(X, y, beta, tau_0, gamma, m):
            X, y = shuffle(X,y)
            for i in range(m):
                tau = tau_0 / (1 + gamma * i)
                beta = beta - tau * gradient(beta, X[i,:], y[i])
            return beta
```

```
In [16]: def SG_minibatch(X, y, beta, tau_0, gamma, m, B = 1):
        # here we do not replace the samples since there would not be enough samples
        # to perform m=150 optimization steps for larger B
        for i in range(m):
            X, y = shuffle(X,y)
            tau = tau_0 / (1 + gamma * i)
            beta = beta - tau * gradient(beta, X[:B:], y[:B])
        return beta
```

```
In [17]: def SG_momentum(X, y, beta, tau_0, gamma, mu, m):
        X, y = shuffle(X,y)
        # initialize g
        g = np.zeros(len(beta))
        for i in range(m):
            tau = tau_0 / (1 + gamma * i)
            g = mu * g + (1 - mu) * gradient(beta, X[i:], y[i])
            beta = beta - tau * g
        return beta
```

```
In [18]: def ADAM(X, y, beta, m, tau=10**-4, mu1=0.9, mu2=0.999, eps=10**-8):
        X, y = shuffle(X,y)
        # initialize g, q
        g = np.zeros(len(beta))
        q = np.zeros(len(beta))
        for i in range(m):
            grad = gradient(beta, X[i:], y[i])
            g = mu1 * g + (1 - mu1) * grad
            q = mu2 * q + (1 - mu2) * np.square(grad)
            beta = beta - tau / (np.sqrt(q) + eps) * g
        return beta
```

```

In [19]: # so wie im Skript
def Newton_Raphson(X, y, beta, m, lambda_=0.001):# hier muss lambda kleiner gewählt w
    N,D = X.shape
    for _ in range(m):
        z = X.dot(beta)
        y_tilde = y / sigmoid(-y * z)
        W = np.diag(lambda_ / N * sigmoid(z) * sigmoid(-z))
        inv = np.linalg.inv(1 / lambda_ * np.identity(D) + np.dot(X.T, np.dot(W, X)))
        beta = beta + np.dot(inv, np.dot(X.T, np.dot(W, y_tilde))) - beta / lambda_
    return beta

In [20]: # so wie auf dem Zettel
def Newton_Raphson1(X, y, beta, m, lambda_=0.001):
    N,D = X.shape
    for _ in range(m):
        z = X.dot(beta)
        y_tilde = y / sigmoid(-y * z)
        W = lambda_ / N * np.diag(sigmoid(z) * sigmoid(-z))
        inv = np.linalg.inv(np.identity(D) + np.dot(X.T, np.dot(W, X)))
        beta = np.dot(inv, np.dot(X.T, W)).dot(z + y_tilde)
    return beta

In [21]: def dual_coordinate_ascent(X, y, m, lambda_=1):
    # initialization:
    N = X.shape[0]
    alpha = np.random.uniform(size=N)
    beta = lambda_ * np.average((alpha * y)[: ,None] * X, axis=0)
    X, y = shuffle(X, y)
    for i in range(m):
        f_prime = y[i] * np.dot(X[i, :], beta) + np.log(alpha[i] / (1 - alpha[i]))
        f_2prime = lambda_ / N * np.dot(X[i, :], X[i, :].T) + 1 / (alpha[i] * (1 - alpha[i]))
        alpha_old = alpha[i].copy()
        alpha[i] = max(0, min(1, alpha[i] - f_prime / f_2prime))
        beta += lambda_ / N * y[i] * X[i, :] * (alpha[i] - alpha_old)
    return beta

In [22]: def SAGD(X, y, beta_init, tau_0, gamma, m, lambda_ = 0.001):
    #init
    try:
        X, y = shuffle(X, y)
        N = X.shape[0]
        beta = beta_init
        g_stored = np.zeros(N)
        g_vec = np.zeros(N)
        for i in range(N):
            #print(sigmoid(-y[i] * np.dot(X[i, :], beta)))
            g_stored[i] = -y[i] * X[i, :] * sigmoid(-y[i] * np.dot(X[i, :], beta))
        g = np.average(g_stored)

```

```

        #iterate
        for i in range(m):
            g_vec[i] = -y[i] * X[i,:] * sigmoid(-y[i] * np.dot(X[i,:],beta))
            g      = g + 1. / N * (g_vec[i] - g_stored[i])
            tau_t = tau_0 / (1 + gamma * i)
            g_stored[i] = g_vec[i]
            beta  = beta * (1 - tau_t / lambda_) - tau_t * g
        return beta
    except:
        raise NotImplementedError('Not yet implemented')

```

dual coordinate ascent

```

In [23]: X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = 0.3 ,random_state=42)
        beta = dual_coordinate_ascent(X_train, y_train, 150)

```

```

        errors = zero_one_loss(predict(beta, X_test), y_test)
        print(errors / len(y_test))

```

0.027777777777778

SAGD

```

In [24]: X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = 0.3 ,random_state=42)
        beta = np.zeros(65)
        tau_0 = 0.001
        gamma = 0.0001
        beta = SAGD(X_train, y_train, beta, tau_0, gamma, 150)

```

```

        errors = zero_one_loss(predict(beta, X_test), y_test)
        print(errors / len(y_test))

```

ValueError

Traceback (most recent call last)

```

<ipython-input-22-028ad3c55fc7> in SAGD(X, y, beta_init, tau_0, gamma, m, lambda_)
    10             #print(sigmoid(-y[i] * np.dot(X[i,:],beta)))
--> 11             g_stored[i] = -y[i] * X[i, :] * sigmoid(-y[i] * np.dot(X[i,:],beta))
    12             g = np.average(g_stored)

```

ValueError: setting an array element with a sequence.

During handling of the above exception, another exception occurred:

NotImplementedError

Traceback (most recent call last)

```
<ipython-input-24-f91d9cc6115d> in <module>()
      3 tau_0 = 0.001
      4 gamma = 0.0001
----> 5 beta = SAGD(X_train, y_train, beta, tau_0, gamma, 150)
      6
      7 errors = zero_one_loss(predict(beta, X_test), y_test)

<ipython-input-22-028ad3c55fc7> in SAGD(X, y, beta_init, tau_0, gamma, m, lambda_)
     20     return beta
     21 except:
----> 22     raise NotImplementedError('Not yet implemented')
```

NotImplementedError: Not yet implemented

GD

```
In [25]: X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = 0.3 ,random_state=42)
        beta = np.zeros(65)
        beta = GD(X_train, y_train, beta, 0.001, 0.0001, 10)

        errors = zero_one_loss(predict(beta, X_test), y_test)
        print(errors / len(y_test))
```

0.0185185185185

Newton Raphson

```
In [26]: X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = 0.3 ,random_state=42)
        beta = np.zeros(65)
        beta = Newton_Raphson(X_train, y_train, beta, 150)

        errors = zero_one_loss(predict(beta, X_test), y_test)
        print(errors / len(y_test))
```

0.0185185185185

```
In [27]: X_test.shape
```

```
Out[27]: (108, 65)
```

SGD w/o replacement

```
In [28]: X_train, X_test, y_train, y_test = train_test_split (X, y, test_size = 0.3 ,random_state=0)
        beta = np.zeros(65)
        tau_0 = 0.001
        gamma = 0.0001
        beta = SGD(X_train, y_train, beta, tau_0, gamma, 150)

        errors = zero_one_loss(predict(beta, X_test), y_test)
        print(errors / len(y_test))

0.0185185185185
```

SGD with minibatch

```
In [29]: X_train, X_test, y_train, y_test = train_test_split (X, y, test_size = 0.3 ,random_state=0)
        beta = np.zeros(65)
        tau_0 = 0.001
        gamma = 0.0001
        beta = SG_minibatch(X_train, y_train, beta, tau_0, gamma, 150, B = 20)

        errors = zero_one_loss(predict(beta, X_test), y_test)
        print(errors / len(y_test))

0.0185185185185
```

ADAM

```
In [30]: X_train, X_test, y_train, y_test = train_test_split (X, y, test_size = 0.3 ,random_state=0)
        beta = np.zeros(65)
        tau_0 = 0.001
        gamma = 0.0001
        mu = 0.1
        beta = ADAM(X_train, y_train, beta, 150)

        errors = zero_one_loss(predict(beta, X_test), y_test)
        print(errors / len(y_test))

0.0185185185185
```

4 1.3 Comparison

```
In [31]: from sklearn.cross_validation import KFold
        from operator import itemgetter
        import itertools
```

```
/usr/local/lib/python3.6/site-packages/sklearn/cross_validation.py:41: DeprecationWarning: This
    "This module will be removed in 0.20.", DeprecationWarning)
```

```
In [32]: optimizers = ['GD', 'SGD', 'SGD_mini', 'SGD_mom', 'ADAM', 'DualCA', 'NR']
        #tbd: 'SAGD'
```

```
class LogisticRegression:
    def __init__(self, X, y, type_, beta_0):
        self.X = X
        self.y = y
        self.type = type_
        self.beta = beta_0

    def compare_parameters(self, tau_0_, mu_, gamma_, iter_ = 150, num_results_shown = 10):
        if num_results_shown is None:
            num_results_shown = len(tau_0_)*len(mu_)*len(gamma_)
        if self.type in ['NR', 'GD']:
            iter_ = 10
        X_tr, X_test, y_tr, y_test = train_test_split(self.X, self.y, test_size = 0.1)
        cross_val_results = []
        beta = self.beta
        for (tau_0, mu, gamma) in itertools.product(tau_0_, mu_, gamma_):
            kf = KFold(y_tr.shape[0], n_folds=10)
            errors = 0
            for train_idx, val_idx in kf:
                if self.type == 'SGD':
                    beta = SGD(X_tr[train_idx], y_tr[train_idx], beta, tau_0, gamma, iter_)
                elif self.type == 'GD':
                    beta = GD(X_tr[train_idx], y_tr[train_idx], beta, tau_0, gamma, iter_)
                elif self.type == 'SGD_mini':
                    #print('SGD minibatch running with B = ', B)
                    beta = SG_minibatch(X_tr[train_idx], y_tr[train_idx], beta, tau_0, gamma, iter_)
                elif self.type == 'SGD_mom':
                    beta = SG_momentum(X_tr[train_idx], y_tr[train_idx], beta, tau_0, gamma, iter_)
                elif self.type == 'ADAM':
                    #ADAM(X, y, beta, m, tau=10**-4, mu1=0.9, mu2=0.999, eps=10**-8):
                    beta = ADAM(X_tr[train_idx], y_tr[train_idx], beta, iter_)
                elif self.type == 'DualCA':
                    beta = dual_coordinate_ascent(X_tr[train_idx], y_tr[train_idx], beta, iter_)
                elif self.type == 'NR':
                    beta = Newton_Raphson(X_tr[train_idx], y_tr[train_idx], beta, iter_)
                elif self.type == 'SAGD':
                    #beta = self.type(X_tr[train_idx], y_tr[train_idx], beta, tau_0, gamma, iter_)
                    raise NotImplementedError('tbd')

            errors += zero_one_loss(predict(beta, X_tr[val_idx]), y_tr[val_idx])
```



```

        cross_val_results.append((tau_0, mu, gamma, errors))

    sort = np.array(sorted(cross_val_results, key = itemgetter(3)))
    print('\nUsing ' , self.type, ' optimizer.')
    print('Listing results starting with lowest error score: \n')
    display(pd.DataFrame(
        data = sort[:num_results_shown - 1, :3],
        index = sort[:num_results_shown - 1, -1],
        columns = ['tau_0', 'mu', 'gamma'])
        .rename_axis('errors', axis = 'columns'))

    return sort[0]
    #print(sorted(cross_val_results, key = itemgetter(3)))

In [33]: tau, mu, gamma = [0.001, 0.01, 0.1], [0.1, 0.2, 0.5], [0.0001, 0.001, 0.01]
        LR = LogisticRegression(X, y, 'SGD', np.zeros(65))
        LR.compare_parameters(tau, mu, gamma)

```

Using SGD optimizer.

Listing results starting with lowest error score:

errors	tau_0	mu	gamma
0.0000	0.0010	0.1000	0.0100
0.0000	0.0010	0.5000	0.0100
1.0000	0.0010	0.2000	0.0100
2.0000	0.0010	0.2000	0.0010
3.0000	0.0010	0.5000	0.0010
4.0000	0.0010	0.1000	0.0001
4.0000	0.0010	0.1000	0.0010
4.0000	0.0010	0.2000	0.0001
5.0000	0.0010	0.5000	0.0001
11.0000	0.0100	0.5000	0.0100
16.0000	0.0100	0.2000	0.0100
18.0000	0.0100	0.2000	0.0010
21.0000	0.0100	0.1000	0.0100
27.0000	0.0100	0.5000	0.0001
28.0000	0.0100	0.1000	0.0010
38.0000	0.0100	0.5000	0.0010
40.0000	0.1000	0.5000	0.0100
47.0000	0.0100	0.2000	0.0001
62.0000	0.0100	0.1000	0.0001
65.0000	0.1000	0.2000	0.0100
66.0000	0.1000	0.5000	0.0001
73.0000	0.1000	0.1000	0.0100
78.0000	0.1000	0.2000	0.0010

```
78.0000 0.1000 0.5000 0.0010
87.0000 0.1000 0.2000 0.0001
96.0000 0.1000 0.1000 0.0010
```

```
Out[33]: array([ 0.001,  0.1   ,  0.01  ,  0.   ])
```

listing the n best results of each optimizer

```
In [34]: def list_n_best(X, y, n, opti_list, tau, mu, gamma):
    overview = np.empty([7], dtype=[('opt', 'U10'), ('tau', 'f4'), ('mu', 'f4'), ('gamma', 'f4')])
    for i, opt in enumerate(opti_list):
        LR = LogisticRegression(X, y, opt, np.zeros(65))
        best = LR.compare_parameters(tau, mu, gamma, num_results_shown=n)
        overview['opt'][i] = opt
        overview['tau'][i] = best[0]
        overview['mu'][i] = best[1]
        overview['gamma'][i] = best[2]
    print("\n\n Overall best parameters: \n")
    display(pd.DataFrame(
        data = overview, index=overview['opt'],
        columns = ['tau', 'mu', 'gamma']).rename_axis('optimizer', axis = 'columns'))
    return overview
```

```
In [35]: params = list_n_best(X, y, 10, optimizers, tau, mu, gamma)
```

Using GD optimizer.

Listing results starting with lowest error score:

```
errors tau_0      mu gamma
3.0000 0.0010 0.1000 0.0100
3.0000 0.0010 0.2000 0.0001
3.0000 0.0010 0.2000 0.0010
3.0000 0.0010 0.2000 0.0100
3.0000 0.0010 0.5000 0.0001
3.0000 0.0010 0.5000 0.0010
3.0000 0.0010 0.5000 0.0100
3.0000 0.0100 0.1000 0.0001
3.0000 0.0100 0.1000 0.0010
```

Using SGD optimizer.

Listing results starting with lowest error score:

errors	tau_0	mu	gamma
0.0000	0.0010	0.5000	0.0001
1.0000	0.0010	0.2000	0.0010
1.0000	0.0010	0.5000	0.0100
2.0000	0.0010	0.1000	0.0010
2.0000	0.0010	0.1000	0.0100
2.0000	0.0010	0.2000	0.0100
5.0000	0.0010	0.2000	0.0001
6.0000	0.0010	0.1000	0.0001
6.0000	0.0010	0.5000	0.0010

Using SGD_mini optimizer.
Listing results starting with lowest error score:

errors	tau_0	mu	gamma
1.0000	0.0010	0.2000	0.0001
2.0000	0.0010	0.1000	0.0010
2.0000	0.0010	0.5000	0.0010
2.0000	0.0010	0.5000	0.0100
3.0000	0.0010	0.1000	0.0001
3.0000	0.0010	0.1000	0.0100
3.0000	0.0010	0.2000	0.0010
3.0000	0.0010	0.2000	0.0100
3.0000	0.0010	0.5000	0.0001

Using SGD_mom optimizer.
Listing results starting with lowest error score:

errors	tau_0	mu	gamma
1.0000	0.0010	0.1000	0.0100
1.0000	0.0010	0.2000	0.0010
1.0000	0.0010	0.2000	0.0100
1.0000	0.0010	0.5000	0.0100
2.0000	0.0010	0.2000	0.0001
2.0000	0.0010	0.5000	0.0010
3.0000	0.0010	0.1000	0.0010
4.0000	0.0010	0.5000	0.0001
4.0000	0.0100	0.2000	0.0100

Using ADAM optimizer.

Listing results starting with lowest error score:

errors	tau_0	mu	gamma
1.0000	0.0100	0.1000	0.0100
1.0000	0.1000	0.5000	0.0001
2.0000	0.0010	0.2000	0.0001
2.0000	0.0010	0.2000	0.0010
2.0000	0.0100	0.2000	0.0010
2.0000	0.0100	0.2000	0.0100
2.0000	0.0100	0.5000	0.0001
2.0000	0.1000	0.1000	0.0001
2.0000	0.1000	0.2000	0.0010

Using DualCA optimizer.

Listing results starting with lowest error score:

errors	tau_0	mu	gamma
14.0000	0.0010	0.1000	0.0001
16.0000	0.0100	0.2000	0.0100
16.0000	0.0100	0.5000	0.0100
17.0000	0.1000	0.5000	0.0001
18.0000	0.1000	0.2000	0.0001
19.0000	0.0100	0.1000	0.0001
19.0000	0.0100	0.5000	0.0001
19.0000	0.1000	0.2000	0.0100
21.0000	0.0010	0.1000	0.0100

Using NR optimizer.

Listing results starting with lowest error score:

errors	tau_0	mu	gamma
27.0000	0.0010	0.1000	0.0001
27.0000	0.0010	0.1000	0.0010
27.0000	0.0010	0.1000	0.0100
27.0000	0.0010	0.2000	0.0001
27.0000	0.0010	0.2000	0.0010
27.0000	0.0010	0.2000	0.0100
27.0000	0.0010	0.5000	0.0001
27.0000	0.0010	0.5000	0.0010

27.0000 0.0010 0.5000 0.0100

Overall best parameters:

optimizer	tau	mu	gamma
GD	0.0010	0.1000	0.0100
SGD	0.0010	0.5000	0.0001
SGD_mini	0.0010	0.2000	0.0001
SGD_mom	0.0010	0.1000	0.0100
ADAM	0.0100	0.1000	0.0100
DualCA	0.0010	0.1000	0.0001
NR	0.0010	0.1000	0.0001

speed test

```
In [105]: import matplotlib.pyplot as plt
          %matplotlib inline

data = np.empty(7, dtype=[ ('opt', 'U10'),
                           ('tau', 'f4'),
                           ('mu', 'f4'),
                           ('gamma', 'f4'),
                           ('compC', 'i8'),
                           ('iterdata', '(8, 4)f4')]) #use for iterations-training
#do iterations from 1 = 2^0 to 256=2^8 in steps of x2 as standard, ergo 9 steps. for
# iterations from 1 to 17 in steps of +2

for i, struct in enumerate(data):
    struct['opt'] = params['opt'][i]
    struct['tau'] = params['tau'][i]
    struct['mu'] = params['mu'][i]
    struct['gamma'] = params['gamma'][i]
    if 0:#struct['opt'] in ['NR', 'GD']:
        struct['iterdata'] = [[1.,0.,0.,0.],
                               [5.,0.,0.,0.],
                               [9.,0.,0.,0.],
                               [13.,0.,0.,0.],
                               [17.,0.,0.,0.],
                               [21.,0.,0.,0.],
                               [25.,0.,0.,0.],
                               [29.,0.,0.,0.]]
    else:
        struct['iterdata'] = [[1.,0.,0.,0.],
```

```

[2.,0.,0.,0.],
[4.,0.,0.,0.],
[8.,0.,0.,0.],
[16.,0.,0.,0.],
[32.,0.,0.,0.],
[64.,0.,0.,0.],
[128.,0.,0.,0.]]

#insert computational complexiies
D = 65
N = 357
B = 10
data['compC'] = D
data['compC'][data['opt'] == 'GD'] = N*D
data['compC'][data['opt'] == 'SGD_mini'] = B*D
data['compC'][data['opt'] == 'NR'] = N*D**2

X_tr, X_test, y_tr, y_test = train_test_split (X, y, test_size = 0.3 ,random_state =
for i, struct in enumerate(data):
    beta = np.zeros(65)
    tau_0 = struct['tau']
    mu = struct['mu']
    gamma = struct['gamma']
    for j, iterdata in enumerate(struct['iterdata']):
        iter_ = int(iterdata[0])
        errors = 0
        if struct['opt'] == 'SGD':
            beta = SGD(X_tr, y_tr, beta, tau_0, gamma, iter_)
        elif struct['opt'] == 'GD':
            beta = GD(X_tr, y_tr, beta, tau_0, gamma, iter_)
        elif struct['opt'] == 'SGD_mini':
            #print('SGD minibatch running with B = ', B)
            beta = SG_minibatch(X_tr, y_tr, beta, tau_0, gamma, iter_, B = 10)
        elif struct['opt'] == 'SGD_mom':
            beta = SG_momentum(X_tr, y_tr, beta, tau_0, gamma, mu, iter_)
        elif struct['opt'] == 'ADAM':
            #ADAM(X, y, beta, m, tau=10**-4, mu1=0.9, mu2=0.999, eps=10**-8):
            beta = ADAM(X_tr, y_tr, beta, iter_)
        elif struct['opt'] == 'DualCA':
            beta = dual_coordinate_ascent(X_tr, y_tr, iter_)
        elif struct['opt'] == 'NR':
            beta = Newton_Raphson(X_tr, y_tr, beta, iter_)
        elif struct['opt'] == 'SAGD':
            #beta = self.type(X_tr, y_tr, beta, tau_0, gamma, mu, iter_)
            raise NotImplementedError('tbd')

```

```

        iterdata[1] = zero_one_loss(predict(beta, X_tr), y_tr)/len(y_tr)
        iterdata[2] = zero_one_loss(predict(beta, X_test), y_test)/len(y_test)
        iterdata[3] = iterdata[0]*struct['compC']
        #print(struct['iterdata'][:,3])
        #print(struct['iterdata'])
    plt.figure()
    plt.semilogx(struct['iterdata'][:,3],struct['iterdata'][:,1],label='training')
    plt.semilogx(struct['iterdata'][:,3],struct['iterdata'][:,2],label='test_'+s)
    plt.ylabel('error')
    plt.xlabel('estimated T')
    plt.legend()
    #plt.plot(iterdata[3],iterdata[2],label='test')

```







